

EDITH — Project Review Document

Even Dead, I'm The Hero

A Browser-Native AI Agent for Real-Time Web Automation

1. Project Overview

EDITH is an **AI-powered browser extension** that allows users to automate real-world browser tasks through natural language. A user simply types "Go to YouTube and search for VS Code tutorials" and EDITH autonomously opens YouTube, finds the search box, types the query, presses Enter, and reports back — all inside the user's actual browser, with full visual feedback.

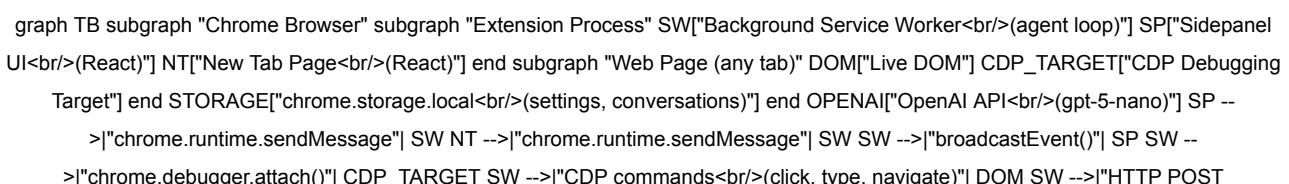
It is **not** a chatbot with a web wrapper. It is not a headless scraper. It is an autonomous agent that sees, reasons, and acts inside a live Chrome tab, exactly as a human would.

2. Technology Stack

Layer	Technology	Why This Choice
Extension Runtime	WXT (https://wxt.dev/) (Web Extension Toolkit)	Framework-agnostic extension build tool. Handles manifest generation, hot reload, entrypoint routing, and cross-browser support. Compiles TypeScript → Chrome MV3 extension.
Language	TypeScript	Type safety across the entire codebase. Catches bugs at compile time — critical for an agent that sends CDP commands.
UI Framework	React 18	Sidepanel chat UI. React's state model handles real-time progress updates, conversation history, and mode switching cleanly.
Browser Automation	Chrome DevTools Protocol (CDP) via chrome.debugger API	The core innovation — explained in detail below.
LLM Interface	OpenAI SDK (openai npm package)	Any OpenAI-compatible API (gpt-4o-mini, gpt-5-nano, local LLMs via LMStudio/Ollama).
Storage	chrome.storage.local	All data stays on-device. No backend server, no database.
Build & Bundle	Vite 7 (via WXT)	Sub-4-second production builds. Tree-shaking removes dead code.

3. Architecture — Deep Dive

3.1 System Architecture Diagram



/chat/completions" | OPENAI SW -->|"read/write"| STORAGE SP -->|"read"| STORAGE style SW fill:#5c73f2,color:#fff style OPENAI
fill:#10a37f,color:#fff style STORAGE fill:#333,color:#fff

3.2 The Three Execution Contexts

Chrome extensions run code in **three isolated contexts**. Understanding this is critical:

Context	File	Runs Where	Lifespan	Can Access
Background Service Worker	entrypoints/background/index.ts	Extension's V8 isolate	Persistent while extension is loaded	chrome.* APIs, network, storage. Cannot access DOM.
Sidepanel UI	entrypoints/sidepanel/App.tsx	Rendered in Chrome's side panel	While panel is open	React rendering, chrome.runtime.sendMessage to talk to background
Content/Injected Script	Injected via Runtime.evaluate	Inside the target web page's JS context	Per-page execution	Full DOM access — document.querySelector, window, etc.

The service worker **orchestrates everything**. It receives user messages from the sidepanel, calls the LLM, and sends CDP commands to the target tab. It never touches the DOM directly — that happens via injected scripts.

4. Browser Automation — The Core Innovation

4.1 Why CDP via `chrome.debugger`?

Most browser automation tools (Selenium, Playwright, Puppeteer) use **headless Chrome** — a separate browser process with no visible window. This has three fatal problems:

1. **Websites detect headless Chrome** — missing `window.chrome, navigator.webdriver === true`, missing GPU fingerprint
2. **User's real session isn't available** — no cookies, no login state, no saved passwords
3. **Requires installing a separate binary** — Chromium, chromedriver, or a Node.js server

EDITH solves all three by using the `chrome.debugger` API:

```
// Attach to any open tab - the user's REAL Chrome session
await chrome.debugger.attach({ tabId: 12345 }, '1.3');

// Send CDP commands to that tab
await chrome.debugger.sendCommand(
  { tabId: 12345 },
  'Input.dispatchMouseEvent',
  { type: 'mousePressed', x: 400, y: 300, button: 'left', clickCount: 1 }
);
```

This is the **same protocol** that Chrome DevTools uses. When you open Chrome DevTools and click "Elements" or "Console", it uses CDP internally. Our extension does the same thing, but programmatically.

4.2 The Automation Pipeline

```

sequenceDiagram participant U as User participant SP as Sidepanel participant BG as Service Worker participant LLM as OpenAI API
participant TAB as Browser Tab U->>SP: "Search YouTube for VS Code tutorials" SP->>BG: { type: "AGENT_RUN", prompt: "..." } BG-->
>>SP: ack (immediately) loop Agent Loop (max 30 steps) BG->>LLM: System prompt + history + tools LLM-->>BG: tool_call:
open_browser("https://youtube.com") BG->>TAB: chrome.tabs.create({ url }) BG->>TAB: chrome.debugger.attach() BG-->>SP: progress("□
open_browser") BG->>LLM: tool result + history LLM-->>BG: tool_call: take_snapshot() BG->>TAB: Runtime.evaluate(snapshotScript) TAB-->>BG: { elements: [{uid:1, tag:"input", ...}, ...] } BG->>LLM: snapshot data LLM-->>BG: tool_call: type_text(uid:1, "vs code tutorials") BG-->>TAB: Input.dispatchKeyEvent (per character) BG->>LLM: result LLM-->>BG: tool_call: press_key("Enter") BG->>TAB: Input.dispatchKeyEvent("Enter") BG->>LLM: result LLM-->>BG: tool_call: task_complete("Searched YouTube for VS Code tutorials") BG-->>TAB: chrome.debugger.detach() BG-->>SP: { type: "agent_done" } end SP->>U: "□ Searched YouTube for VS Code tutorials"

```

4.3 DOM Snapshot System

The agent needs to "see" the page. Instead of screenshots (expensive, requires vision models), we **inject a JavaScript function** into the page that walks the DOM tree and extracts all interactive elements:

```

const snapshotScript = function() {
    const elements = [];
    const walker = document.createTreeWalker(document.body, NodeFilter.SHOW_ELEMENT);

    while (walker.nextNode()) {
        const node = walker.currentNode;
        // Only collect interactive elements: inputs, buttons, links
        if (isClickable(node) || hasRole(node)) {
            elements.push({
                uid: counter++,           // Stable reference ID
                tag: 'input',
                name: node.getAttribute('aria-label') || node.innerText,
                placeholder: node.getAttribute('placeholder'),
                x: rect.left, y: rect.top,   // For click coordinates
                isInput: true,
                isClickable: true,
            });
        }
    }
    return JSON.stringify({ url: location.href, title: document.title, elements });
};

```

The LLM receives output like:

```

URL: https://www.youtube.com
Title: YouTube

ELEMENTS (uid | type | label):
1 | INPUT | "Search"
2 | BUTTON | "Search"
3 | LINK | "Home"
14 | VIDEO | "VS Code Tutorial for Beginners"

```

It then calls `type_text(uid=1, text="vs code tutorials")` and the service worker resolves UID 1 → coordinates (x=400, y=52) → sends a CDP mouse click at that point → then dispatches key events for each character.

4.4 CDP Typing — Character-Level Input

Most web apps (YouTube, Google, React SPAs) use JavaScript event listeners, not native HTML form handling. Simply setting `input.value = "text"` does **not** work — React never sees the change.

Our approach dispatches **real keyboard events** through CDP:

```
for (const char of text) {  
    // 1. keyDown - tells the browser a key was pressed (no text insertion)  
    await cdp(tabId, 'Input.dispatchKeyEvent', {  
        type: 'keyDown',  
        key: char,  
        code: `Key${char.toUpperCase()}`,  
        windowsVirtualKeyCode: char.toUpperCase().charCodeAt(0),  
        // NOTE: No `text` field here - avoids double-insertion  
    });  
  
    // 2. char - this is what actually inserts the character  
    await cdp(tabId, 'Input.dispatchKeyEvent', {  
        type: 'char',  
        key: char,  
        text: char,  
        unmodifiedText: char,  
    });  
  
    // 3. keyUp - release the key  
    await cdp(tabId, 'Input.dispatchKeyEvent', {  
        type: 'keyUp',  
        key: char,  
    });  
  
    await sleep(40); // Human-like typing speed  
}
```

This three-phase `keyDown → char → keyUp` cycle is exactly what happens when a human presses a key. React's synthetic event system picks it up correctly.

5. LLM Integration

5.1 Model-Agnostic Design

EDITH doesn't hardcode any model. It uses the OpenAI Chat Completions API format — any provider that speaks this protocol works:

Provider	Model	Cost	Works?
OpenAI	gpt-4o-mini, gpt-5-nano	Pay-per-use	<input type="checkbox"/>
Google	Gemini (via OpenAI-compatible endpoint)	Free tier available	<input type="checkbox"/>
Ollama	llama3, CodeLlama (local)	Free	<input type="checkbox"/>
LMStudio	Any GGUF model (local)	Free	<input type="checkbox"/>
Groq	llama3-70b	Free tier	<input type="checkbox"/>

The user just changes the **API Base URL** and **Model** in settings. No code changes needed.

5.2 Tool Calling (Function Calling)

The LLM doesn't execute actions directly. Instead, we define **tools** using OpenAI's function calling schema:

```
{  
  name: 'click',  
  description: 'Click an element by its UID from the snapshot.',  
  parameters: {  
    type: 'object',  
    properties: {  
      uid: { type: 'number', description: 'UID from snapshot' },  
    },  
    required: ['uid'],  
  },  
}
```

The LLM returns structured JSON:

```
{  
  "tool_calls": [{  
    "id": "call_abc123",  
    "function": {  
      "name": "click",  
      "arguments": "{\"uid\": 14}"  
    }  
  }]  
}
```

Our service worker parses this, executes the action, and feeds the result back. This loop continues until the LLM calls `task_complete()`.

5.3 Tools Available

Tool	Purpose
<code>open_browser(url)</code>	Open a new tab with a URL
<code>navigate(url)</code>	Navigate the current tab
<code>take_snapshot()</code>	Get all interactive elements with UIDs
<code>click(uid)</code>	Click an element by UID
<code>type_text(uid, text)</code>	Type text into an input field
<code>press_key(key)</code>	Press Enter, Tab, Escape, etc.
<code>scroll(direction, amount)</code>	Scroll the page
<code>screenshot()</code>	Capture the visible page
<code>task_complete(summary)</code>	Signal task is done, release browser

5.4 Context Window Management

A browsing task can take 10-20 steps. Each step generates snapshot data (potentially hundreds of elements). If we send the entire history every time, we:

- Exceed the model's context window

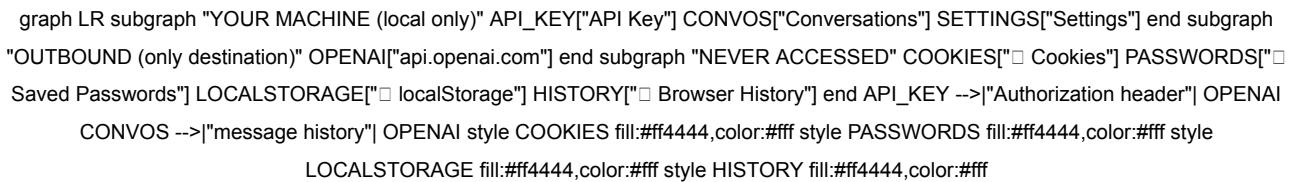
- Waste tokens (= money)
- Trigger safety filters from accumulated page data

Solution: `pruneHistory()` — before every LLM call, we keep only:

- All user messages (so the model remembers the original task)
- The last 6 rounds of tool interactions
- This keeps context under ~8K tokens per call

6. Security Architecture

6.1 Data Flow — Nothing Leaves Without You Knowing



6.2 Security Measures Implemented

Measure	Code Location	What It Prevents
Password fields excluded from snapshot	automation.ts line 175	Credentials never reach LLM
Hidden fields excluded from snapshot	automation.ts line 175	Hidden tokens never exposed
Input .value never read	automation.ts line 190	Typed content stays local
No fetch() calls in codebase	Verified via grep	No data exfiltration to third parties
No document.cookie access	Verified via grep	Session hijacking impossible
No chrome.cookies API	Not in manifest.json	Extension cannot read any cookies
chrome.storage.local only	storage.ts	Data never synced to Chrome cloud
Debugger auto-detaches on completion	background/index.ts	Extension releases tab control after task

7. Extension Manifest (Permissions Explained)

```
{
  "manifest_version": 3,
  "permissions": [
    "debugger",           // CDP access to tabs for automation
    "sidePanel",          // Show chat UI in Chrome's side panel
    "storage",             // Save settings & conversations locally
    "tabs",                // Query which tab is active
    "activeTab",           // Access the current tab's URL/title
    "scripting",            // Inject snapshot script into pages
    "alarms",               // Scheduled tasks (future feature)
    "notifications"        // Notify user when tasks complete (future)
  ],
  "host_permissions": ["<all_urls>"] // Automate any website
}
```

Every permission is justified. `host_permissions: <all_urls>` is required because the agent needs to automate **any** site the user asks (YouTube, Amazon, GitHub, etc.).

8. Message Protocol

Communication between sidebar and service worker uses `chrome.runtime.sendMessage`:

Chat Mode (synchronous):

```
Sidepanel → { type: "CHAT", prompt: "What is React?" }
Background → calls OpenAI → waits → responds
Background → { ok: true, conversationId: "abc-123" }
Sidepanel → reads updated conversation from storage
```

Agent Mode (fire-and-forget + events):

```
Sidepanel → { type: "AGENT_RUN", prompt: "Go to YouTube..." }
Background → { ok: true } (immediate ack)
Background → starts agent loop...
  → broadcasts: { type: "agent_progress", text: "❑ open_browser" }
  → broadcasts: { type: "agent_progress", text: "❑ Snapshot: YouTube" }
  → broadcasts: { type: "agent_progress", text: "❑ type_text" }
  → broadcasts: { type: "agent_done", conversationId: "abc-123" }
Sidepanel → reads final conversation from storage
```

The fire-and-forget pattern is critical because agent tasks can take 30-90 seconds. Chrome's `sendMessage` has a timeout — if we awaited the full response, it would fail silently.

9. File Structure

```

extension/
├── wxt.config.ts          # WXT build configuration
├── package.json            # Dependencies
└── tsconfig.json           # TypeScript configuration

|
├── entrypoints/
|   ├── background/
|   |   └── index.ts          # Service worker - agent loop, message handler
|   ├── sidebar/
|   |   ├── index.html        # Sidebar HTML shell
|   |   └── main.tsx          # React mount point
|   └── newtab/
|       ├── index.html        # New tab HTML shell
|       └── main.tsx          # Clock + quick task input

|
├── lib/
|   ├── automation.ts        # CDP browser control (450 lines)
|   ├── llm.ts                # OpenAI API client
|   ├── agent.ts              # System prompt, tools, snapshot formatter
|   └── storage.ts            # chrome.storage.local CRUD

|
├── assets/
|   └── global.css            # EDITH design system (CSS variables)

|
└── .output/
    └── chrome-mv3/           # Built extension (load this in Chrome)
        ├── manifest.json
        ├── background.js
        ├── sidebar.html
        ├── newtab.html
        └── chunks/
            └── assets/

```

10. Anticipated Panel Questions & Answers

Q1: "Why not use Selenium or Playwright instead of CDP?"

Selenium and Playwright launch **separate headless browser processes**. This means:

- The user's cookies, login sessions, and saved passwords are not available
- Websites easily detect headless Chrome (`navigator.webdriver === true`)
- It requires installing Chromium binaries (700MB+)

CDP via `chrome.debugger` operates **inside the user's actual Chrome browser**. The browser fingerprint is real. The session is real. The cookies are real. No additional software is needed — it's a 280KB extension.

Q2: "How does the LLM know what's on the page?"

We inject a DOM traversal script via `Runtime.evaluate` that walks the full DOM tree and extracts:

- Every interactive element (inputs, buttons, links, videos)
- Their semantic labels (`aria-label`, `title`, `placeholder`, `innerText`)
- Their bounding box coordinates (for click targeting)
- A unique integer UID for reference

This creates a **structured text representation** of the page. The LLM reads text like "`UID=3 | INPUT | Search`" and responds with `type_text(uid=3, text="vs code")`. It never sees raw HTML, screenshots, or the visual layout — just a clean element list.

Q3: "What if the LLM calls the wrong tool or clicks the wrong element?"

The agent loop has multiple safeguards:

1. **Max 30 steps** — prevents infinite loops
2. `task_complete() tool` — LLM explicitly signals when done; debugger detaches immediately
3. **Snapshot invalidation** — after every click/type action, the old snapshot is discarded. The agent must call `take_snapshot()` again to see the updated page.
4. **Error handling** — if a CDP command fails, the error is sent back to the LLM as a tool result, allowing it to try an alternative approach

However, the agent IS imperfect. LLMs sometimes:

- Click the wrong UID
- Navigate to an unrelated site (mitigated by system prompt restrictions)
- Loop unnecessarily

This is a fundamental limitation of LLM-based agents — no AI agent today is 100% reliable.

Q4: "Is this extension safe to use on personal accounts?"

Yes, with caveats:

- **Chrome/Google:** No risk. `chrome.debugger` is an official API used by thousands of extensions.
- **YouTube/Amazon:** No risk for personal browsing tasks. These platforms target bots doing mass scraping, ad fraud, or automated purchases — not personal automation.
- **LinkedIn:** Low risk for browsing. Do NOT automate bulk messaging or connection requests — LinkedIn actively bans that.
- **Data safety:** No cookies, passwords, or session tokens leave the machine. Only visible page labels are sent to OpenAI.

Q5: "Why didn't you use browser vision (screenshots + GPT-4V) instead of text snapshots?"

Three reasons:

1. **Cost:** GPT-4V with vision costs ~10-50x more per call than text-only models. A 10-step task with vision could cost \$0.50-\$2.
2. **Speed:** Encoding and transmitting screenshot images adds 2-5 seconds per step.
3. **Accuracy:** Text-based element lists with UIDs give the LLM **exact coordinates** for clicking. Vision models guess pixel coordinates and are often wrong by 20-50px.

Text snapshots are faster, cheaper, and more accurate for structured web pages.

Q6: "How does typing work in React applications like YouTube search?"

React uses a virtual DOM and synthetic event system. Setting `input.value = "text"` directly does **not** trigger React's `onChange` handler. The input visually updates but React's internal state stays empty.

Our solution dispatches CDP `Input.dispatchKeyEvent` calls that generate **real browser keyboard events**:

- `keyDown` (key press signal — no text insertion)
- `char` (actual character insertion)
- `keyUp` (key release)

These are indistinguishable from a human typing. React's event delegation picks them up through `document.addEventListener`, and `onChange` fires correctly.

We initially had a bug where `keyDown` carried the `text` field, causing each character to be inserted twice ("vvss ccooddee"). The fix was ensuring only the `char` event carries the `text` payload.

Q7: "What happens if the service worker goes to sleep?"

Chrome MV3 service workers have a 5-minute idle timeout. However, our agent loop maintains active network connections (LLM API calls) and CDP commands, which keep the service worker alive. Chrome doesn't terminate a service worker that has active `chrome.debugger` connections.

If the worker does terminate mid-task (e.g., Chrome update), the debugger auto-detaches and the task stops. The conversation history is already saved to storage, so the user can see how far it got.

Q8: "Can this extension be published to the Chrome Web Store?"

Technically yes, but with restrictions:

- Extensions using the `debugger` permission require extra review by Google
 - The `host_permissions: <all_urls>` requires justification
 - Google may request a detailed privacy policy explaining what data goes to OpenAI
 - Many automation-focused extensions exist on the store (Selenium IDE, Automa, etc.) — there is precedent
-

Q9: "How is this different from BrowserOS, the open-source AI browser?"

Aspect	BrowserOS	EDITH
Architecture	Custom Chromium fork + extension	Standard Chrome extension only
Browser requirement	Must download their custom browser	Works in any Chrome/Edge
Automation API	Custom <code>chrome.browserOS.*</code> (proprietary)	Standard <code>chrome.debugger</code> (CDP)
LLM provider	Multi-provider via Vercel AI SDK	Any OpenAI-compatible API
User's real session	<input type="checkbox"/> Separate browser = separate cookies	<input type="checkbox"/> Same Chrome = same session
Install size	~700MB (full Chromium)	~280KB (extension only)
MCP support	<input type="checkbox"/> Built-in	<input type="checkbox"/> Phase 2

EDITH's key advantage is **zero friction** — no new browser install, everything works in the user's existing Chrome with their existing logins.

Q10: "What are the ethical implications of an AI that controls a browser?"

Valid concern. Guardrails implemented:

1. **User must initiate every task** — the agent never acts autonomously
2. **Chrome shows a yellow banner** — "An extension is debugging this browser" is always visible
3. **task_complete() detaches debugger** — agent physically cannot interact after finishing
4. **System prompt restricts navigation** — blocks wandering to unrelated sites
5. **No purchase automation** — the system prompt says "never confirm purchases without explicit user approval"

The extension is a tool, not an autonomous system. It cannot install itself, cannot run without the user's Chrome, and cannot bypass the debugger permission dialog.