

EDITH — Complete Technical Architecture Report

EDITH — *Elite Digital Intelligence & Task Handler* A Chrome Extension-based AI agent with autonomous browser automation, multi-tab parallel research, and intelligent document processing.

Table of Contents

1. [Technology Stack](#)
 2. [High-Level System Architecture](#)
 3. [Extension Architecture — The Core Agent](#)
 4. [Modes of Operation](#)
 5. [Data Flow](#)
 6. [Browser Automation Engine](#)
 7. [Multi-Tab Research Orchestrator](#)
 8. [Tool Ecosystem](#)
 9. [Document Intelligence Module](#)
 10. [Storage & State Management](#)
 11. [LLM Integration](#)
 12. [Project File Structure](#)
 13. [Security Architecture](#)
-

1. Technology Stack

1.1 Chrome Extension — Core Agent

Technology	Version	Purpose
WXT	0.20.17	Chrome extension framework (Manifest V3)
TypeScript	5.7.3	Type-safe logic across all modules

Technology	Version	Purpose
React	18.3.1	Sidepanel & New Tab UI
React DOM	18.3.1	DOM rendering for extension pages
OpenAI SDK	4.86.2	LLM API calls directly from browser context
TailwindCSS	3.4.17	Utility-first styling for extension UI
PostCSS	8.5.3	CSS processing pipeline
Chrome Debugger API	CDP	DOM automation via Chrome DevTools Protocol

Chrome Extension Permissions: debugger · sidePanel · storage · tabs · activeTab · scripting · alarms · notifications · <all_urls>

1.2 Document Intelligence Module


Technology	Purpose
FastAPI + Uvicorn	REST API server
Sentence-Transformers (intfloat/e5-large-v2)	Text embedding model
FAISS (CPU)	Vector similarity search
PyMuPDF (fitz)	PDF parsing with hyperlink extraction
python-docx	Word document parsing
python-pptx	PowerPoint parsing + embedded image OCR
OpenPyXL + Pandas	Excel/CSV parsing & analysis
Pillow + Pytesseract	OCR for images and embedded visuals
Google GenAI (gemini-2.5-flash)	Primary LLM for Q&A
OpenAI (gpt-5-nano)	Fallback LLM
LangChain + LangGraph	Interactive reasoning agent orchestration
BeautifulSoup4	Web scraping for linked documents

1.3 External AI Services

Service	Model	Role
OpenAI	gpt-5-nano (configurable)	Primary agent LLM – reasoning, tool-calling, intent
OpenAI	gpt-4	Interactive reasoning agent (Document Intelligence)
Google Gemini	gemini-2.5-flash	Primary LLM for document Q&A (with key rotation)

2. High-Level System Architecture

EDITH is composed of **two independent systems** that together form the project:



```

API["FastAPI Server<br/>POST /api/v1/hackrx/run"]
DP["Document Parser<br/>(PDF/DOCX/PPTX/XLSX/IMG)"]
EMB["Embedding Engine<br/>(e5-large-v2)"]
FAISS["FAISS Vector Index"]
RAG["RAG Pipeline<br/>(Gemini / OpenAI)"]
AGENT["Interactive Agent<br/>(LangGraph + GPT-4)"]

API --> DP --> EMB --> FAISS
API --> RAG
FAISS --> RAG
RAG -->|"if interactive"| AGENT
end

LLM_EXT -->|"OpenAI API"| OAI["OpenAI<br/>gpt-5-nano"]

style SP fill:#6C5CE7,color:#fff
style BG fill:#FDCB6E,color:#000
style AM fill:#00B894,color:#fff
style RES fill:#74B9FF,color:#fff
style API fill:#E17055,color:#fff
style RAG fill:#A29BFE,color:#fff

```

3. Extension Architecture – The Core Agent


The extension follows a **3-layer architecture**: UI → Background Service Worker → Automation Layer.

3.1 Layer Breakdown

```

----- ENTRYPOINTS
|
|
| Sidepanel (entrypoints/sidepanel/)      New Tab
(entrypoints/newtab/) |
|   main.tsx  → React mount             |   main.tsx
|   |
|   └─ App.tsx   → Chat UI, settings    |   └─ index.html
|
|
|
|   — Communicates via chrome.runtime.sendMessage() —

```



```
| agent.ts           → System prompt, browser tool definitions, snapshot  
fmt |  
| automation.ts    → CDP-based DOM automation (1186 lines)  
|  
| research.ts      → Multi-tab research decompose/execute/aggregate  
|  
| llm.ts           → OpenAI SDK wrapper (callLLM, message formatting)  
|  
| storage.ts       → chrome.storage.local CRUD for  
settings/conversations |  
| tab_manager.ts   → Tab lifecycle: create, detach, track active tabs
```

3.2 Key Source Files & Responsibilities

File	Lines	Role
background/index.ts	715	Service worker: message routing, agent loop, research runner, alarm handler
agent.ts	398	System prompt (140 lines), BROWSER_TOOLS definitions (15 tools), formatSnapshot() , pruneHistory() .
automation.ts	1186	CDP automation: takeSnapshot() (JS injection), clickElement() , typeText() , pressKey() , scrollPage() , selectOption() , hoverElement() , setvalue()

File	Lines	Role
<u>research.ts</u>	523	<u>decomposeTask()</u> , <u>runSubTask()</u> , <u>aggregateResults()</u> – MapReduce-style parallel research
<u>llm.ts</u>	99	OpenAI SDK client, message format conversion, tool schema mapping
<u>storage.ts</u>	134	Typed CRUD for settings, conversations (100 max), MCP servers, scheduled tasks
<u>tab_manager.ts</u>	~100	<u>createTab()</u> , <u>detachAll()</u> , multi-tab lifecycle for research

4. Modes of Operation

Every user message enters the background service worker and is routed to one of **four execution modes**:

4.1 CHAT Mode – Simple Conversation

```
sequenceDiagram
    participant U as User (Sidepanel)
    participant BG as Background Worker
    participant LLM as OpenAI API

    U->>BG: { type: "CHAT", prompt: "Hello!" }
    BG-->>BG: Load settings + conversation
    BG-->>LLM: callLLM(prompt, history, tools[])
    Note over BG,LLM: No browser tools provided<br/>Pure text
    conversation
    LLM-->>BG: "Hi! I can help you browse..."
    BG-->>BG: Save to chrome.storage.local
    BG-->>U: sendResponse({ ok: true })
```

Characteristics: - **No browser tools** are passed to the LLM (empty tools array) - Synchronous response — sidepanel waits for sendResponse - Conversation persisted to chrome.storage.local - Used for: greetings, questions about capabilities, general chat

4.2 AGENT_RUN Mode — Single-Tab Browser Automation

sequenceDiagram

```
participant U as User (Sidepanel)
participant BG as Background Worker
participant LLM as OpenAI API
participant CDP as Chrome Debugger
participant TAB as Browser Tab
```

U->>BG: { type: "AGENT_RUN", prompt: "Search Amazon for laptop" }

BG-->>U: { ok: true } (immediate ack)

Note over U: Sidepanel now listens
for broadcastEvent()

loop Agent Loop (max 30 steps)

BG->>LLM: callLLM(SYSTEM_PROMPT, history, BROWSER_TOOLS)

alt LLM returns tool_calls

```
    LLM-->>BG: [{ name: "open_browser", args: {url: "amazon.com"} }]  
        BG-->>BG: progress("👉 open_browser: ...")
```

BG-->>CDP: openBrowser("https://amazon.com")

CDP-->>TAB: chrome.debugger.attach + Page.navigate

CDP-->>BG: tabId

BG-->>BG: sleep(1500ms)

Note over BG: Next iteration...

else LLM returns tool: take_snapshot

BG-->>CDP: takeSnapshot(tabId)

CDP-->>TAB: Runtime.evaluate(SNAPSHOT_JS)

TAB-->>CDP: { url, title, elements[], rawText }

CDP-->>BG: PageSnapshot

BG-->>BG: formatSnapshot(snapshot)

Note over BG: Elements formatted with

UUIDs
Prioritized: inputs → buttons → links

else LLM returns tool: click(uid)

BG-->>CDP: clickElement(uid, snapshot, tabId)

CDP-->>TAB: Input.dispatchEvent(x, y)

BG-->>BG: sleep(1200ms)

BG-->>CDP: auto-takeSnapshot()

```

        Note over BG: Auto-snapshot after every action
    else LLM returns task_complete
        LLM-->>BG: { name: "task_complete", args: {summary:
    "Done!"} }

        BG->>BG: Save final message
        BG->>CDP: detachDebugger(tabId)
        BG->>U: broadcastEvent({ type: "agent_done" })
        Note over BG: EXIT agent loop
    else LLM returns text (no tools)
        LLM-->>BG: "Task completed."
        BG->>U: broadcastEvent({ type: "agent_done" })
        Note over BG: EXIT agent loop
    end
end

```

Characteristics: - **Asynchronous** — sends immediate ack, then broadcasts progress events - **Max 30 steps** per run - **15 browser tools** provided via `BROWSER_TOOLS` from `agent.ts` - **Auto-snapshot** after every mutating action (click, type_text, press_key, hover, select_option, set_value) - **Snapshot loop detection** — warns after 3 consecutive snapshots without action - **New tab detection** — automatically switches `activeTabId` when clicks open new tabs - **Abort support** — `agentAbortFlag` checked before each iteration and each tool execution

4.3 RESEARCH_RUN Mode — Multi-Tab Parallel Research

```

sequenceDiagram
    participant U as User (Sidepanel)
    participant BG as Background worker
    participant LLM as OpenAI API
    participant TM as Tab Manager
    participant T1 as Tab 1
    participant T2 as Tab 2
    participant T3 as Tab 3

    U->>BG: { type: "RESEARCH_RUN", prompt: "Compare iPhone 16 prices"
}

    BG-->>U: { ok: true } (immediate ack)

    rect rgb(240, 248, 255)
        Note over BG,LLM: Phase 1: Task Decomposition
        BG->>LLM: decomposeTask(prompt)
        LLM-->>BG: { isResearch: true, subTasks: [<br/> {url:

```

```

"amazon.com", goal: "iPhone 16 price"},<br/> {url: "flipkart.com",
goal: "iPhone 16 price"},<br/> {url: "apple.com", goal: "official
price"}<br/>] }

end

rect rgb(255, 248, 240)

    Note over BG,T3: Phase 2: Parallel Execution
    BG->>TM: createTab(amazon.com) → tabId1
    BG->>TM: createTab(flipkart.com) → tabId2
    BG->>TM: createTab(apple.com) → tabId3
    BG->>BG: sleep(2000ms)

par Run all tabs simultaneously
    BG->>T1: runSubTask(settings, subTask1, tabId1)
    Note over T1: Agent loop: snapshot → type → search →
extract
    and
    BG->>T2: runSubTask(settings, subTask2, tabId2)
    Note over T2: Agent loop: snapshot → type → search →
extract
    and
    BG->>T3: runSubTask(settings, subTask3, tabId3)
    Note over T3: Agent loop: snapshot → type → search →
extract
end

rect rgb(240, 255, 240)

    Note over BG,LLM: Phase 3: Result Aggregation
    BG->>LLM: aggregateResults(prompt, subTaskResults[])
    LLM-->>BG: "iPhone 16 prices:\n• Amazon: ₹79,900\n• Flipkart:
₹78,999\n• Apple: ₹79,900"
end

    BG->>TM: detachAll()
    BG->>U: broadcastEvent({ type: "agent_done" })

```

Characteristics:

- **MapReduce pattern** — decompose → parallel execute → synthesize
- **Max 5 concurrent tabs** (configurable via `MAX_RESEARCH_TABS`) - Each sub-task tab runs its own independent agent loop with `extract_data` as the terminal tool (instead of `task_complete`)
- **Minimum 2 sub-tasks** required — otherwise falls back to Agent mode
- Uses `Promise.allSettled()` — partial failures don't kill the entire research
- Tabs remain open after research so the user can manually review sources

4.4 AGENT_STOP Mode – Emergency Halt

```
sequenceDiagram
    participant U as User
    participant BG as Background Worker
    participant LOOP as Agent/Research Loop

    U->>BG: { type: "AGENT_STOP" }
    BG-->>BG: agentAbortFlag = true
    BG-->>U: { ok: true }

    Note over LOOP: Next iteration checks flag
    LOOP-->>LOOP: if (agentAbortFlag) break
    LOOP-->>BG: detachDebugger()
    LOOP-->>U: broadcastEvent("agent_done")
```

Characteristics: - Sets a global agentAbortFlag boolean - Checked at the **top of each agent iteration** and **before each tool execution** - Gracefully saves conversation state before stopping - Detaches CDP debugger to free browser resources

5. Data Flow

5.1 End-to-End Agent Flow

```
flowchart TD
    A[" User types in Sidepanel"] -->
    B["chrome.runtime.sendMessage()"]
    B --> C{"Message Router<br/>(background/index.ts)"}
    C -->|"type: CHAT"| D["handlechat()"]
    C -->|"type: AGENT_RUN"| E["runAgent()"]
    C -->|"type: RESEARCH_RUN"| F["runResearchFromPrompt()"]
    C -->|"type: AGENT_STOP"| G["Set abort flag"]

    D --> H["callLLM(prompt, [], no tools)"]
    H --> I["Save response → chrome.storage"]
    I --> J["sendResponse to sidepanel"]

    E --> K["callLLM(SYSTEM_PROMPT,<br/>pruneHistory(messages),<br/>BROWSER_TOOLS)"]
```

```

K --> L{"LLM Response"}
L -->| "tool_calls" | M["Execute via CDP"]
M --> N["Auto-snapshot"]
N --> O["Append result to history"]
O --> K
L -->| "task_complete" | P["Save + broadcast done"]
L -->| "text only" | P

F --> Q["decomposeTask(prompt)"]
Q --> R["tabManager.createTab() × N"]
R --> S["Promise.allSettled(<br/>runSubTask per tab)"]
S --> T["aggregateResults(results)"]
T --> P

style C fill:#FDCB6E,color:#000
style M fill:#00B894,color:#fff
style K fill:#6C5CE7,color:#fff
style Q fill:#74B9FF,color:#fff

```

5.2 Snapshot Lifecycle

The snapshot is the **central data structure** that bridges the LLM and the browser:



```

end

subgraph SnapshotElement
    C3
    E1["uid: number"]
    E2["tag, role, name"]
    E3["x, y, width, height"]
    E4["context (parent section)"]
    E5["value, placeholder"]
    E6["checked, disabled, options[]"]
end


```

Snapshot processing pipeline: 1. **Inject** `SNAPSHOT_JS` into the page via `CDP Runtime.evaluate` 2. **Scan** all DOM elements — filter by clickable tags, actionable ARIA roles, and input elements 3. **Classify** each element's type: `INPUT`, `BUTTON`, `LINK`, `CHECKBOX`, `SELECT`, `RADIO`, etc. 4. **Assign UIDs** — sequential numeric IDs based on position 5. **Extract context** — walk up to 5 parent levels for `aria-label`, `id`, `class` hints 6. **Prioritize** — Tier 1 (inputs) → Tier 2 (buttons/checkboxes) → Tier 3 (product cards) → Tier 4 (links) 7. **Format** as a flat text string with `uid | TYPE | "label" [in: section-context]` 8. **Append** the page's `rawText` (first 5000 chars) for context

6. Browser Automation Engine (CDP)

The automation layer in `automation.ts` uses the **Chrome DevTools Protocol** via `chrome.debugger`:

6.1 CDP Architecture



```

F["Page.navigate"]
G["Runtime.evaluate"]
H["Input.dispatchMouseEvent"]
I["Input.dispatchKeyEvent"]
J["DOM.getDocument"]
K["Page.captureScreenshot"]
L["Network events"]

end

A --> B --> D
A --> C --> E --> F & G & H & I & J & K & L

```

6.2 Key Automation Functions

Function	CDP Method	What It Does
openBrowser(url)	chrome.tabs.create + <u>ensureAttached</u>	Opens new tab, attaches debugger, navigates
<u>navigateTo(url)</u>	Page.navigate	In-tab navigation
takeSnapshot()	Runtime.evaluate(SNAPSHOT_JS)	Injects 300-line JS scan interactive DC elements
clickElement(uid)	Input.dispatchMouseEvent	Maps UID → (x,y) from snapshot, dispatches mouse events
typeText(text, uid)	Runtime.evaluate(focus) + Input.dispatchKeyEvent	Character-by- character typing wi 30-80ms delays
pressKey(key)	Input.dispatchKeyEvent (keyDown + keyUp)	Simulates Enter, T Escape, arrows
scrollPage(dir)	Runtime.evaluate (window.scrollBy)	Scroll up/down/top/botto
selectOption(uid, value)	Runtime.evaluate	Programmatically sets select value + fires change event
hoverElement(uid)	Input.dispatchMouseEvent (mouseMoved)	Triggers hover state for dropdown menu
setValue(uid, value)	Runtime.evaluate	Direct .value injection for numer

Function	CDP Method	What It Does
		fields
takeScreenshot()	Page.captureScreenshot	PNG screenshot of visible viewport
waitForNetworkIdle()	Network event monitoring	Waits until no pending network requests
<u>waitForLoad()</u>	Page.loadEventFired	Waits for page load completion
detachDebugger()	chrome.debugger.detach	Releases debugger from tab

6.3 Multi-Tab Support

The automation layer tracks attached tabs via `attachedTabs: Set<number>: -`
 Multiple debugger sessions can be active simultaneously (for research) -
`lastSingleTabId` provides backward compatibility for single-tab agent runs -
`detachAllDebuggers()` cleans up all sessions at once

7. Multi-Tab Research Orchestrator

Defined in `research.ts`, the research system uses a **3-phase MapReduce** architecture:

7.1 Phase 1: Task Decomposition (decomposeTask)

An LLM call analyzes the user's prompt and decides: - **Is this a research task?** (`isResearch: boolean`) - **What sub-tasks are needed?** Each with: `description, url, extractionGoal`

The LLM is given strict criteria for what counts as research: > Genuine research requiring **2+ different websites** – NOT simple browsing, messaging, or single-site tasks.

7.2 Phase 2: Parallel Execution (runSubTask)

Each sub-task gets: - Its own Chrome tab via `tabManager.createTab(url)` - Its own agent loop (similar to `runAgent` but with `extract_data` as terminal tool) - Its own LLM conversation context - An abort signal connection to the global abort flag

Sub-task tools are the same as BROWSER_TOOLS but with task_complete replaced by extract_data (which captures findings instead of ending the run).

7.3 Phase 3: Aggregation (aggregateResults)

An LLM call receives all extracted data and the original prompt, then synthesizes:

- Key findings from each source - Differences and similarities across sources
 - Source citations for every data point - Tables or bullet points for comparisons
-

8. Tool Ecosystem

8.1 Browser Tools (defined in agent.ts)

Tool	Description	Auto-Snapshots?
task_complete	Signals task completion — stops the agent immediately	—
open_browser	Opens URL in new tab	No (manual snapshot needed)
take_snapshot	Gets all interactive elements with UIDs	—
click	Clicks element by UID	<input checked="" type="checkbox"/> Auto-snapshot after
type_text	Types text char-by-char into field by UID	<input checked="" type="checkbox"/> Auto-snapshot after
press_key	Keyboard events (Enter, Tab, Escape, etc.)	<input checked="" type="checkbox"/> Auto-snapshot after
scroll	Scrolls page up/down	No
screenshot	Captures visual screenshot	No
select_option	Selects dropdown option by value	<input checked="" type="checkbox"/> Auto-snapshot after
hover	Triggers hover/mouseover on element	<input checked="" type="checkbox"/> Auto-snapshot after
set_value	Direct value injection for inputs	<input checked="" type="checkbox"/> Auto-snapshot after
wait_for_page_update	Waits for AJAX/network idle	<input checked="" type="checkbox"/> Auto-snapshot after

Tool	Description	Auto-Snapshots?
<u>navigate</u>	In-tab URL navigation	No


8.2 Research-Only Tools

Tool	Description
extract_data	Terminal tool for research sub-tasks — captures extracted information

9. Document Intelligence Module

A standalone FastAPI microservice for intelligent document Q&A, located at `Document Intelligence/doc_intel/`.

9.1 Architecture



```

O --> P["build_faiss_index(embeddings)"]
P --> Q["save_to_cache()<br/>(pickle + FAISS)"]

Q --> R["ThreadPoolExecutor<br/>(8-16 workers)"]
R --> S["process_question() × N"]

S --> T["FAISS search: top-15 chunks"]
T --> U{"Contains URLs/API links?"}

U -->| "No" | V["Standard RAG<br/>TEMPLATE.format(clauses, question)
<br/>→ Gemini 2.5 Flash"]
U -->| "Yes" | W["Interactive Agent<br/>(LangGraph + GPT-4)"]

V --> X["Gemini fails?"]
X -->| "Yes" | Y["Fallback: OpenAI gpt-5-nano"]
X -->| "No" | Z["Return answer"]
Y --> Z
W --> Z

style A fill:#E17055,color:#fff
style P fill:#00CEC9,color:#fff
style V fill:#6C5CE7,color:#fff
style W fill:#FDCB6E,color:#000

```

9.2 Dual-Path Processing

The system dynamically chooses between two processing paths:

Path A: Standard RAG Pipeline

For regular documents without interactive instructions: 1. Parse document → chunks 2. Embed with e5-large-v2 → FAISS index 3. For each question: semantic search → top-15 chunks → LLM prompt 4. Gemini 2.5 Flash (primary) → OpenAI (fallback)

Path B: Interactive Reasoning Agent

For documents containing URLs, API endpoints, or multi-step instructions: 1. Detects that relevant chunks contain URLs/API references via `contains_api_or_url()` 2. Activates the **LangGraph ReAct agent** with two tools: - `document_retriever` — searches the loaded document via FAISS - `web_scraper_tool` — fetches data from external URLs/APIs 3. The agent reasons through multi-step instructions autonomously

9.3 Key Features

Feature	Implementation
Multi-format support	PDF, DOCX, PPTX, XLSX/XLS, JPG/PNG/GIF (OCR)
Linked document expansion	Auto-extracts URLs from document text, fetches and indexes linked content
Embedding caching	MD5-hashed filenames in pdf_cache/ — persistent across requests
Concurrent Q&A	ThreadPoolExecutor with 8-16 workers for batch question processing
LLM key rotation	3 Gemini API keys cycled via itertools.cycle()
OneDrive/SharePoint	Auto-converts sharing URLs to direct download URLs
OCR for images in PPTX	Extracts embedded images from slides and runs Pytesseract
Prompt engineering	Domain-specific template for insurance, legal, HR documents

9.4 Document Intelligence File Structure

```
Document Intelligence/doc_intel/
├── app/
│   ├── main.py                  # FastAPI app, token auth, request
│   └── logging
|       ├── document_parser.py    # parse_pdf, parse_docx, parse_pptx,
|       │   parse_excel, parse_image
|       ├── embeddings.py         # SentenceTransformer (e5-large-v2) +
|       │   FAISS index builder
|       ├── retrieval.py          # RAG pipeline, caching, linked doc
|       │   expansion, question processing
|       ├── interactive_agent.py   # LangGraph ReAct agent with
|       │   document_retriever + web_scraper tools
|       ├── prompt_template.py    # Domain-specific RAG prompt template
|       └── utils.py              # clean_response(),
|           contains_api_or_url()
└── pdf_cache/                  # Persistent embedding cache (pickle +
    FAISS)
└── requirements.txt             # 23 Python dependencies
```

```
|── Dockerfile          # Docker containerization support  
└── .env                # Gemini keys (x3), OpenAI key, Groq  
key
```

10. Storage & State Management

10.1 Extension Storage (chrome.storage.local)

```
graph LR  
    subgraph "chrome.storage.local"  
        A["edith_api_key<br/>(OpenAI API key)"]  
        B["edith_api_base_url<br/>(default: api.openai.com/v1)"]  
        C["edith_model<br/>(default: gpt-4o-mini)"]  
        D["edith_conversations<br/>(last 100 conversations)"]  
        E["edith_mcp_servers<br/>(MCP server configs)"]  
        F["edith_schedules<br/>(scheduled tasks)"]  
    end
```

10.2 Data Structures

```
interface Conversation {  
    id: string;           // UUID  
    title: string;         // First 60 chars of prompt  
    messages: Message[];   // Full message history  
    createdAt: number;  
    updatedAt: number;  
}  
  
interface Message {  
    id: string;  
    role: 'user' | 'assistant' | 'tool';  
    content: string;  
    toolCalls?: ToolCall[]; // For assistant messages with tool  
                           // calls  
    toolCallId?: string;    // For tool response messages  
    toolName?: string;  
    timestamp: number;  
}  
  
interface ScheduledTask {  
    id: string;  
    name: string;
```

```

prompt: string;           // Agent prompt to execute
cronExpression: string;
enabled: boolean;
lastRun?: number;
}

```

10.3 Document Intelligence Caching

- **In-memory cache:** pdf_cache dict (keyed by document URL)
 - **Persistent cache:** pdf_cache/ directory with pickled pages + embeddings
 - **Cache key:** MD5 hash of the document URL
 - **FAISS index:** Rebuilt from cached embeddings on load (fast)
-

11. LLM Integration

11.1 Extension LLM Client

The extension calls OpenAI-compatible APIs directly from the browser via the **OpenAI SDK**:

```

const client = new OpenAI({
  apiKey: settings.apiKey,
  baseURL: settings.apiUrl,           // User-configurable
  dangerouslyAllowBrowser: true,       // Required for extension
  context
});

const response = await client.chat.completions.create({
  model: settings.model,             // User-configurable (default:
  gpt-4o-mini)
  messages: [system, ...history],
  tools: browserTools,
  tool_choice: 'auto',
  max_completion_tokens: 4096,
});

```

11.2 System Prompt Architecture

The SYSTEM_PROMPT in agent.ts (140 lines) contains:

- **Workflow instructions:** open → snapshot → interact → verify → complete
- **Snapshot reading guide:** How to interpret UIDs, types, and section context
- **Site-specific intelligence:** YouTube, WhatsApp, Gmail, Amazon, e-commerce

patterns - **Forbidden actions:** Never re-open sites, never click navigation links when searching - **Completion criteria:** URL-based verification for each task type

11.3 Conversation History Pruning

`pruneHistory()` keeps only the **last 6 tool exchange rounds** to manage token costs:
- User/assistant text messages are always kept
- Only the most recent tool call/response pairs are retained

11.4 Document Intelligence LLM Strategy

Question → FAISS top-15 chunks → prompt template

Primary: Gemini 2.5 Flash (3 API keys rotated)

↓ on failure

Fallback: OpenAI gpt-5-nano

↓ on failure

Return: "Unable to process this query"

If interactive (URLs/APIs detected):

→ LangGraph Agent with OpenAI GPT-4

12. Project File Structure

```
EDITH - Final yr project/
|
└── EDITH-main/EDITH-main/
    |
    └── extension/                               ★ CORE AGENT
        |
        └── wxt.config.ts                      # Manifest v3 config +
    permissions
        |
        └── package.json                         # Dependencies (OpenAI,
    React, WXT, TS)
        |
        └── tsconfig.json                        # TypeScript configuration
        |
        └── tailwind.config.js                  # TailwindCSS theme
        |
        └── postcss.config.js                   # PostCSS plugins
        |
        └── entrypoints/
            |
            └── background/
                |
                └── index.ts                     # ★ Service worker (715
                                                lines)
                |
                └── |                                # Message router, agent
                    |
                    └── loop,
```

```
|           |           |           # research runner, alarm  
handler  
|           |           |           sidepanel/  
|           |           |           main.tsx      # React mount  
|           |           |           App.tsx       # Chat UI + settings  
|           |           |           index.html  
|           |           newtab/  
|           |           |           main.tsx      # New tab page  
|           |           |           index.html  
|           |           lib/  
|           |           |           agent.ts      # ★ System prompt (140  
lines) +  
|           |           |           # BROWSER_TOOLS  
definitions +  
|           |           |           # formatSnapshot() +  
pruneHistory()  
|           |           |           automation.ts # ★ CDP automation (1186  
lines)  
|           |           |           # Snapshot JS, click,  
type, scroll,  
|           |           |           # hover, select, multi-  
tab attach  
|           |           |           research.ts   # ★ Research orchestrator  
(523 lines)  
|           |           |           # decomposeTask,  
runSubTask,  
|           |           |           # aggregateResults  
|           |           |           llm.ts        # OpenAI SDK wrapper (99  
lines)  
|           |           |           storage.ts    # chrome.storage CRUD (134  
lines)  
|           |           |           tab_manager.ts # Multi-tab lifecycle  
(~100 lines)  
|           |           assets/  
|           |           |           global.css  
|           |           Document Intelligence/  
|           |           doc_intel/  
|           |           |           app/  
|           |           |           main.py      # FastAPI app + auth  
middleware  
|           |           document_parser.py # PDF/DOCX/PPTX/XLSX/Image  
parsers
```

```

    |   └── embeddings.py
builder
    |   └── retrieval.py
linked docs
    |   └── interactive_agent.py
    |   └── prompt_template.py
template
    |   └── utils.py
    └── pdf_cache/
cache
    ├── requirements.txt
    ├── Dockerfile
    └── .env
OpenAI, Groq

```

13. Security Architecture

Layer	Mechanism	Details
Extension Permissions	Manifest V3	Scoped: debugger, sidePanel, storage, tabs, scripting, alarms
API Key Storage	chrome.storage.local	User's OpenAI key stored locally, never transmitted to any server except OpenAI
LLM Calls	Direct browser → OpenAI	No intermediary – extension calls API directly with dangerouslyAllowBrowser: true
Agent Safety	Abort flag	User can immediately stop any running agent via AGENT_STOP
Snapshot Loop Prevention	Counter + warning	Detects 3+ consecutive snapshots without action
New Tab Tracking	activeTabId switch	Automatically follows cross-tab navigation
CDP Cleanup	finally block	detachDebugger() always called, even on errors

Layer	Mechanism	Details
Doc Intel Auth	Bearer token	Static token validation on all API requests
Gemini Key Rotation	<code>itertools.cycle()</code>	3 keys rotated to avoid rate limits
Conversation Limits	100 max conversations	Prevents unbounded storage growth
History Pruning	Last 6 tool rounds	Prevents token overflow on long agent runs

Report generated from exhaustive source code analysis of every file in the extension and Document Intelligence modules. All architecture diagrams and data flows are derived directly from the actual implementation.